

CarND-P2-Report Traffic Sign Classification

1. Dataset Summary and Exploration

1.1. Basic Summary of Dataset

Number of training examples = 34799

Number of validation examples = 4410

Number of testing examples = 12630

Image data shape = (32, 32)

Number of classes = 43

1.2. Explorative Visualization of Dataset

In cell 120 of Jupyter notebook, we can visualize the frequency by Traffic sign ID in training dataset.




Figure 1. Frequency (Percentage) by Traffic Sign ID in Training Dataset

In cell 121-122, I imported the "signnames.csv" and found one example for each traffic sign ID.




Figure 2. Visualization of Traffic Sign 0

2. Design and Test a Model Architecture

2.1 Data Preprocessing

In cell 123-124, I converted the images from RGB to grayscale using package cv2. This is because the color information may not be useful in this task. Using 1 channel instead of 3 channels may also reduce the computational cost in deep learning neural network.

The following is an example of “End of no passing” before and after grayscaling:




Figure 3. “End of no passing” before and after Grayscale

2.2 Model Architecture

The basic architecture is as follows:

Layer 1: Convolutional. Input 32x32x1, the filter 5x5x1, stride = [1,1,1,1], the output 28x28x6.

Activation: RELU

Pooling: The output 14x14x6.

Layer 2: Convolutional. Input 14x14x6, the filter 5x5x1, stride = [1,1,1,1], the output 10x10x16.

Activation: RELU

Pooling: The output 5x5x16.

Flatten using `tf.contrib.layers.flatten`

Layer 3: Fully Connected. Input 400, the output 120.

Activation: RELU.

Layer 4: Fully Connected. Input 120, output 84.

Activation: RELU

Layer 5: Fully Connected (Logits). This should have 43 outputs.

The cell 126 implemented this architecture.

2.3 Model Training and Results

The training data set includes 34,799 samples,

The validation data set includes 4,410 samples,

The testing data set includes 12,630 samples.

The optimizer we used is “tf.train.AdamOptimizer”

The learning rate is set as 0.001.

Batch size is set as 16.

The epoch number is set as 100.

During training, the best validation accuracy is recorded, so is the corresponding model.

In cell 130, we can see that the best validation accuracy is 0.945.

In cell 86, we can see the test accuracy using the best model is 0.901.

3. Test the Model on New Images

3.1 Visualization of Five New Images



Figure 4. Five New Images on Website.

These five images are randomly selected. The original image is not 32x32, “cv2.resize()” function is applied to change the image size.

The five images are also converted to grayscale images using the same function.

3.2 Performance on the New Images

In cell 104, the best model was applied on the five new images. The test accuracy is 100%.

3.3 Model Certainty – Softmax Probabilities

Cell 108 and 109 calculated the softmax probabilities and printed out the top five probabilities for each image.

Probabilities (each row represents the top five softmax probabilities for each image)

```
[ 1.00000000e+00, 8.25672494e-20, 1.21157593e-36, 0.00000000e+00, 0.00000000e+00],  
[ 1.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00],  
[ 1.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00],  
[ 1.00000000e+00, 1.26217107e-16, 1.24292340e-19, 4.45523719e-21, 1.63005973e-23],  
[ 1.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00]
```

Corresponding Traffic Sign id:

```
[ 2, 1, 3, 0, 4],  
[14, 0, 1, 2, 3],  
[17, 0, 1, 2, 3],  
[28, 30, 11, 1, 25],  
[33, 0, 1, 2, 3]
```


Combining with Figure 4, we can see that the deep neural network labels the correct traffic IDs for the five new images with an absolute high probability 100%.

4. Visualization of Hidden Layers

In order to verify that the CNN really can capture some important features in images, it is interesting to visualize the hidden layers of the CNN.

Recall that in the model architecture of this project, the output of the first convolutional layer is 14x14x6 and the output of the second convolution layer is 5x5x16. For the best model we found, the visualization of first and second convolution layers for the “Speed Limit (50 km/h)” in Figure 4 are as follows:

Visualizing layer of Trained CNN: conv1



Visualizing layer of Trained CNN: conv2

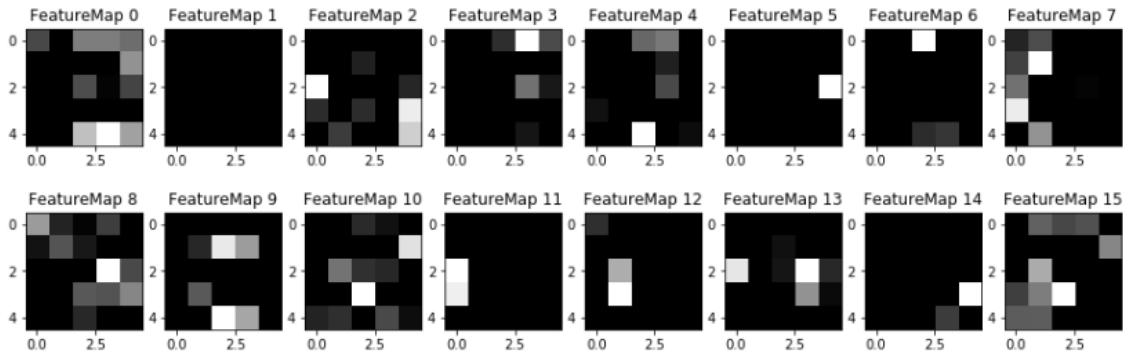



Figure 5. Visualization of Hidden Layers in the Learned Best Model

Similarly, we can visualize the hidden layers of a CNN with random weights for the same input image:

Visualizing layer of Randomized CNN: conv1



Visualizing layer of Randomized CNN: conv2

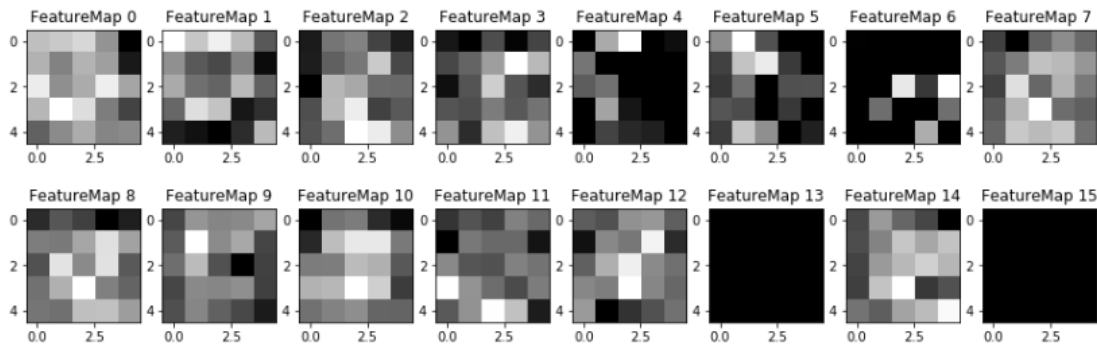


Figure 6. Visualization of Hidden Layers in a CNN with Random Weights

It is obvious that a well-trained CNN model really captures some features comparing with a randomized CNN model.